

Name: KRISHVANTH KUMAR E
REG NO:192125063

Set-I

1. .(i) Write a function called kelvin_to_celsius() that takes a temperature in Kelvin and returns that temperature in Celsius (**Hint:** To convert from Kelvin to Celsius you subtract 273.15) (ii) Write suitable R code to compute the mean, median ,mode of the following values c(90, 50, 70, 80, 70, 60, 20, 30, 80, 90, 20)
(iii) Write R code to find 2nd highest and 3rd Lowest value of above problem.

The screenshot shows an R online compiler interface. The code in the editor is:

```
main.r
1 kelvin_to_celsius <- function(temp_kelvin) {
2   temp_celsius <- temp_kelvin - 273.15
3   return(temp_celsius)
4 }
5 values <- c(90, 50, 70, 80, 70, 60, 20, 30, 80, 90, 20)
6
7 # Mean
8 mean_value <- mean(values)
9 cat("Mean:", mean_value, "\n")
10
11 # Median
12 median_value <- median(values)
13 cat("Median:", median_value, "\n")
14
15 # Mode
16 mode_value <- names(table(values))[table(values) == max(table(values))]
17 cat("Mode:", mode_value, "\n")
18 # 2nd highest value
19 sorted_values <- sort(unique(values), decreasing = TRUE)
20 second_highest <- sorted_values[2]
21 cat("Second highest value:", second_highest, "\n")
22
23 # 3rd lowest value
24 third_lowest <- sorted_values[3]
25 cat("Third lowest value:", third_lowest, "\n")
```

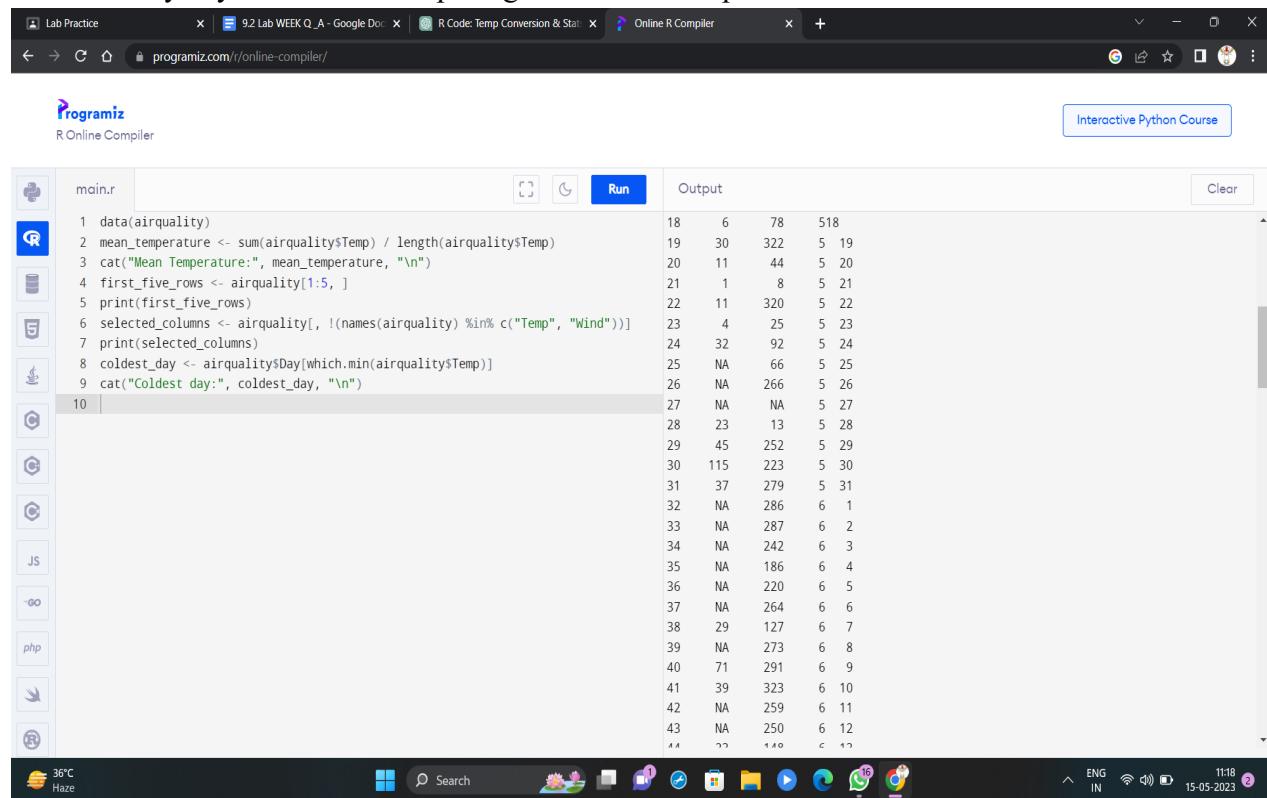
The output window shows the results of the R script execution:

```
Rscript /tmp/ITI55SSvPf.r
Mean: 60
Median: 70
Mode: 20 70 80 90
Second highest value: 80
Third lowest value: 70
```

The Windows taskbar at the bottom of the screenshot includes icons for search, file explorer, and various system status indicators like battery level and network connection.

2. Explore the airquality dataset. It contains daily air quality measurements from New York during a period of five months:
- Ozone: mean ozone concentration (ppb),
 - Solar.R: solar radiation (Langley),

- Wind: average wind speed (mph),
- Temp: maximum daily temperature in degrees Fahrenheit,
- Month: numeric month (May=5, June=6, and so on),
- Day: numeric day of the month (1-31).
 - i. Compute the mean temperature(don't use build in function)
 - ii. Extract the first five rows from airquality.
 - iii. Extract all columns from airquality *except* Temp and Wind
 - iv. Which was the coldest day during the period?
 - v. How many days was the wind speed greater than 17 mph?



The screenshot shows a web browser window with multiple tabs open. The active tab is titled "Lab Practice" and contains R code. To the right of the code editor is an "Output" pane displaying the results of the R script. Below the browser window is a taskbar with various icons and system status information.

```
main.r
1 data(airquality)
2 mean_temperature <- sum(airquality$Temp) / length(airquality$Temp)
3 cat("Mean Temperature:", mean_temperature, "\n")
4 first_five_rows <- airquality[1:5, ]
5 print(first_five_rows)
6 selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))]
7 print(selected_columns)
8 coldest_day <- airquality$Day[which.min(airquality$Temp)]
9 cat("Coldest day:", coldest_day, "\n")
10
```

| | 18 | 6 | 78 | 518 |
|----|-----|-----|----|-----|
| 19 | 30 | 322 | 5 | 19 |
| 20 | 11 | 44 | 5 | 20 |
| 21 | 1 | 8 | 5 | 21 |
| 22 | 11 | 320 | 5 | 22 |
| 23 | 4 | 25 | 5 | 23 |
| 24 | 32 | 92 | 5 | 24 |
| 25 | NA | 66 | 5 | 25 |
| 26 | NA | 266 | 5 | 26 |
| 27 | NA | NA | 5 | 27 |
| 28 | 23 | 13 | 5 | 28 |
| 29 | 45 | 252 | 5 | 29 |
| 30 | 115 | 223 | 5 | 30 |
| 31 | 37 | 279 | 5 | 31 |
| 32 | NA | 286 | 6 | 1 |
| 33 | NA | 287 | 6 | 2 |
| 34 | NA | 242 | 6 | 3 |
| 35 | NA | 186 | 6 | 4 |
| 36 | NA | 220 | 6 | 5 |
| 37 | NA | 264 | 6 | 6 |
| 38 | 29 | 127 | 6 | 7 |
| 39 | NA | 273 | 6 | 8 |
| 40 | 71 | 291 | 6 | 9 |
| 41 | 39 | 323 | 6 | 10 |
| 42 | NA | 259 | 6 | 11 |
| 43 | NA | 250 | 6 | 12 |
| 44 | 22 | 140 | 6 | 13 |

36°C Haze Search 11:18 ENG IN 15-05-2023

Lab Practice | 9.2 Lab WEEK Q_A - Google Doc | R Code: Temp Conversion & Stat | Online R Compiler

programiz.com/online-compiler/

Programiz

R Online Compiler

Interactive Python Course

Clear

main.r

Run

Output

```
1 data(airquality)
2 mean_temperature <- sum(airquality$Temp) / length(airquality$Temp)
3 cat("Mean Temperature:", mean_temperature, "\n")
4 first_five_rows <- airquality[1:5, ]
5 print(first_five_rows)
6 selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))]
7 print(selected_columns)
8 coldest_day <- airquality$Day(which.min(airquality$Temp))
9 cat("Coldest day:", coldest_day, "\n")
```

| | 44 | 23 | 148 | 6 | 13 |
|----|----|-----|-----|----|----|
| 10 | NA | 332 | 6 | 14 | |
| | 45 | NA | 322 | 6 | 15 |
| | 46 | NA | 322 | 6 | 15 |
| | 47 | 21 | 191 | 6 | 16 |
| | 48 | 37 | 284 | 6 | 17 |
| | 49 | 20 | 37 | 6 | 18 |
| | 50 | 12 | 120 | 6 | 19 |
| | 51 | 13 | 137 | 6 | 20 |
| | 52 | NA | 150 | 6 | 21 |
| | 53 | NA | 59 | 6 | 22 |
| | 54 | NA | 91 | 6 | 23 |
| | 55 | NA | 250 | 6 | 24 |
| | 56 | NA | 135 | 6 | 25 |
| | 57 | NA | 127 | 6 | 26 |
| | 58 | NA | 47 | 6 | 27 |
| | 59 | NA | 98 | 6 | 28 |
| | 60 | NA | 31 | 6 | 29 |
| | 61 | NA | 138 | 6 | 30 |
| | 62 | 135 | 269 | 7 | 1 |
| | 63 | 49 | 248 | 7 | 2 |
| | 64 | 32 | 236 | 7 | 3 |
| | 65 | NA | 101 | 7 | 4 |
| | 66 | 64 | 175 | 7 | 5 |
| | 67 | 40 | 314 | 7 | 6 |
| | 68 | 77 | 276 | 7 | 7 |
| | 69 | 97 | 267 | 7 | 8 |

36°C Haze

Search

11:18 15-05-2023

Lab Practice | 9.2 Lab WEEK Q_A - Google Doc | R Code: Temp Conversion & Stat | Online R Compiler

← → C ⌛ 🔒 programiz.com/online-compiler/ ⌂ ⌃ ⌄ ⌅ ⌆ ⌇ ⌈ ⌉ ⌊ ⌋ ⌊ ⌋ ⌊ ⌋ ⌊ ⌋

Programiz R Online Compiler [Interactive Python Course](#)

main.r

Run

Output

Clear

```
1 data(airquality)
2 mean_temperature <- sum(airquality$Temp) / length(airquality$Temp)
3 cat("Mean Temperature:", mean_temperature, "\n")
4 first_five_rows <- airquality[1:5, ]
5 print(first_five_rows)
6 selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))]
7 print(selected_columns)
8 coldest_day <- airquality$Day[which.min(airquality$Temp)]
9 cat("Coldest day:", coldest_day, "\n")
```

| | 10 | 70 | 97 | 272 | 7 | 9 |
|----|---|----|-----|-----|---|----|
| 1 | data(airquality) | 71 | 85 | 175 | 7 | 10 |
| 2 | mean_temperature <- sum(airquality\$Temp) / length(airquality\$Temp) | 72 | NA | 139 | 7 | 11 |
| 3 | cat("Mean Temperature:", mean_temperature, "\n") | 73 | 10 | 264 | 7 | 12 |
| 4 | first_five_rows <- airquality[1:5,] | 74 | 27 | 175 | 7 | 13 |
| 5 | print(first_five_rows) | 75 | NA | 291 | 7 | 14 |
| 6 | selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))] | 76 | 7 | 48 | 7 | 15 |
| 7 | print(selected_columns) | 77 | 48 | 260 | 7 | 16 |
| 8 | coldest_day <- airquality\$Day[which.min(airquality\$Temp)] | 78 | 35 | 274 | 7 | 17 |
| 9 | cat("Coldest day:", coldest_day, "\n") | 79 | 61 | 285 | 7 | 18 |
| 10 | | 80 | 79 | 187 | 7 | 19 |
| | | 81 | 63 | 220 | 7 | 20 |
| | | 82 | 16 | 7 | 7 | 21 |
| | | 83 | NA | 258 | 7 | 22 |
| | | 84 | NA | 295 | 7 | 23 |
| | | 85 | 80 | 294 | 7 | 24 |
| | | 86 | 108 | 223 | 7 | 25 |
| | | 87 | 20 | 81 | 7 | 26 |
| | | 88 | 52 | 82 | 7 | 27 |
| | | 89 | 82 | 213 | 7 | 28 |
| | | 90 | 50 | 275 | 7 | 29 |
| | | 91 | 64 | 253 | 7 | 30 |
| | | 92 | 59 | 254 | 7 | 31 |
| | | 93 | 39 | 83 | 8 | 1 |
| | | 94 | 9 | 24 | 8 | 2 |
| | | 95 | 16 | 77 | 8 | 3 |

36°C Haze

Search

11:18 15-05-2023

Lab Practice | 9.2 Lab WEEK Q_A - Google Doc | R Code: Temp Conversion & Stat | Online R Compiler

programiz.com/online-compiler/

Programiz
R Online Compiler

Interactive Python Course

Clear

main.r

Run

Output

```
1 data(airquality)
2 mean_temperature <- sum(airquality$Temp) / length(airquality$Temp)
3 cat("Mean Temperature:", mean_temperature, "\n")
4 first_five_rows <- airquality[1:5, ]
5 print(first_five_rows)
6 selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))]
7 print(selected_columns)
8 coldest_day <- airquality$Day(which.min(airquality$Temp))
9 cat("Coldest day:", coldest_day, "\n")
```

| | 96 | 78 | NA | 8 | 4 |
|----|-----|-----|-----|---|----|
| 10 | 97 | 35 | NA | 8 | 5 |
| | 98 | 66 | NA | 8 | 6 |
| | 99 | 122 | 255 | 8 | 7 |
| | 100 | 89 | 229 | 8 | 8 |
| | 101 | 110 | 207 | 8 | 9 |
| | 102 | NA | 222 | 8 | 10 |
| | 103 | NA | 137 | 8 | 11 |
| | 104 | 44 | 192 | 8 | 12 |
| | 105 | 28 | 273 | 8 | 13 |
| | 106 | 65 | 157 | 8 | 14 |
| | 107 | NA | 64 | 8 | 15 |
| | 108 | 22 | 71 | 8 | 16 |
| | 109 | 59 | 51 | 8 | 17 |
| | 110 | 23 | 115 | 8 | 18 |
| | 111 | 31 | 244 | 8 | 19 |
| | 112 | 44 | 190 | 8 | 20 |
| | 113 | 21 | 259 | 8 | 21 |
| | 114 | 9 | 36 | 8 | 22 |
| | 115 | NA | 255 | 8 | 23 |
| | 116 | 45 | 212 | 8 | 24 |
| | 117 | 168 | 238 | 8 | 25 |
| | 118 | 73 | 215 | 8 | 26 |
| | 119 | NA | 153 | 8 | 27 |
| | 120 | 76 | 203 | 8 | 28 |
| | 121 | 118 | 225 | 8 | 29 |
| | 122 | 84 | 237 | 8 | 30 |

36°C Haze

Search

11:18 15-05-2023

Lab Practice | 9.2 Lab WEEK Q_A - Google Doc | R Code: Temp Conversion & Stat | Online R Compiler

programiz.com/r/online-compiler/

Programiz
R Online Compiler

Interactive Python Course

Clear

Run

Output

main.r

```
1 data(airquality)
2 mean_temperature <- sum(airquality$Temp) / length(airquality$Temp)
3 cat("Mean Temperature:", mean_temperature, "\n")
4 first_five_rows <- airquality[1:5, ]
5 print(first_five_rows)
6 selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))]
7 print(selected_columns)
8 coldest_day <- airquality$Day[which.min(airquality$Temp)]
9 cat("Coldest day:", coldest_day, "\n")
```

10

| | 123 | 85 | 188 | 8 | 31 |
|-----|-----|-----|-----|----|----|
| 124 | 96 | 167 | 9 | 1 | |
| 125 | 78 | 197 | 9 | 2 | |
| 126 | 73 | 183 | 9 | 3 | |
| 127 | 91 | 189 | 9 | 4 | |
| 128 | 47 | 95 | 9 | 5 | |
| 129 | 32 | 92 | 9 | 6 | |
| 130 | 20 | 252 | 9 | 7 | |
| 131 | 23 | 220 | 9 | 8 | |
| 132 | 21 | 230 | 9 | 9 | |
| 133 | 24 | 259 | 9 | 10 | |
| 134 | 44 | 236 | 9 | 11 | |
| 135 | 21 | 259 | 9 | 12 | |
| 136 | 28 | 238 | 9 | 13 | |
| 137 | 9 | 24 | 9 | 14 | |
| 138 | 13 | 112 | 9 | 15 | |
| 139 | 46 | 237 | 9 | 16 | |
| 140 | 18 | 224 | 9 | 17 | |
| 141 | 13 | 27 | 9 | 18 | |
| 142 | 24 | 238 | 9 | 19 | |
| 143 | 16 | 201 | 9 | 20 | |
| 144 | 13 | 238 | 9 | 21 | |
| 145 | 23 | 14 | 9 | 22 | |
| 146 | 36 | 139 | 9 | 23 | |
| 147 | 7 | 49 | 9 | 24 | |
| 148 | 14 | 20 | 9 | 25 | |
| 149 | 20 | 102 | 9 | 26 | |

36°C Haze

Search

11:18 15-05-2023

The screenshot shows a web-based R compiler interface. The top navigation bar includes tabs for "Lab Practice", "9.2 Lab WEEK Q_A - Google Doc", "R Code: Temp Conversion & Stats", and "Online R Compiler". The URL in the address bar is "programiz.com//online-compiler/". The main area has a "Programiz" logo and a "R Online Compiler" section. On the left, there's a sidebar with icons for various languages: Python, R, C, C++, Java, JavaScript, Go, PHP, and Swift. The code editor window is titled "main.r" and contains the following R script:

```
1 data(airquality)
2 mean_temperature <- sum(airquality$Temp) / length(airquality$Temp)
3 cat("Mean Temperature:", mean_temperature, "\n")
4 first_five_rows <- airquality[1:5, ]
5 print(first_five_rows)
6 selected_columns <- airquality[, !(names(airquality) %in% c("Temp", "Wind"))]
7 print(selected_columns)
8 coldest_day <- airquality$Day[which.min(airquality$Temp)]
9 cat("Coldest day:", coldest_day, "\n")
```

The "Run" button is highlighted in blue. To the right, the "Output" pane displays the results of the R script execution:

| | 129 | 52 | 92 | 9 | 0 |
|-----|-----|-----|----|----|---|
| 130 | 20 | 252 | 9 | 7 | |
| 131 | 23 | 220 | 9 | 8 | |
| 132 | 21 | 230 | 9 | 9 | |
| 133 | 24 | 259 | 9 | 10 | |
| 134 | 44 | 236 | 9 | 11 | |
| 135 | 21 | 259 | 9 | 12 | |
| 136 | 28 | 238 | 9 | 13 | |
| 137 | 9 | 24 | 9 | 14 | |
| 138 | 13 | 112 | 9 | 15 | |
| 139 | 46 | 237 | 9 | 16 | |
| 140 | 18 | 224 | 9 | 17 | |
| 141 | 13 | 27 | 9 | 18 | |
| 142 | 24 | 238 | 9 | 19 | |
| 143 | 16 | 201 | 9 | 20 | |
| 144 | 13 | 238 | 9 | 21 | |
| 145 | 23 | 14 | 9 | 22 | |
| 146 | 36 | 139 | 9 | 23 | |
| 147 | 7 | 49 | 9 | 24 | |
| 148 | 14 | 20 | 9 | 25 | |
| 149 | 30 | 193 | 9 | 26 | |
| 150 | NA | 145 | 9 | 27 | |
| 151 | 14 | 191 | 9 | 28 | |
| 152 | 18 | 131 | 9 | 29 | |
| 153 | 20 | 223 | 9 | 30 | |

The output also includes the message "Coldest day: 5". The status bar at the bottom shows system information like weather (36°C Haze), search, file explorer, taskbar icons, and system status.

3. (i) Get the Summary Statistics of air quality dataset
- (ii) Melt airquality data set and display as a long – format data?
- (iii) Melt airquality data and specify month and day to be “ID variables”?
- (iv) Cast the molten airquality data set with respect to month and date features
- (v) Use cast function appropriately and compute the average of Ozone, Solar.R , Wind and temperature per month?

The screenshot shows the Programiz Online R Compiler interface. On the left, there's a sidebar with various icons for file operations like Open, Save, Print, and Run. The main area has tabs for 'Lab Practice', '9.2 Lab WEEK Q_A - Google Doc', 'R Code: Temp Conversion & Stat...', and 'Online R Compiler'. The current tab is 'Online R Compiler'. A URL 'programiz.com/r/online-compiler/' is visible in the address bar. The code editor contains the following R script:

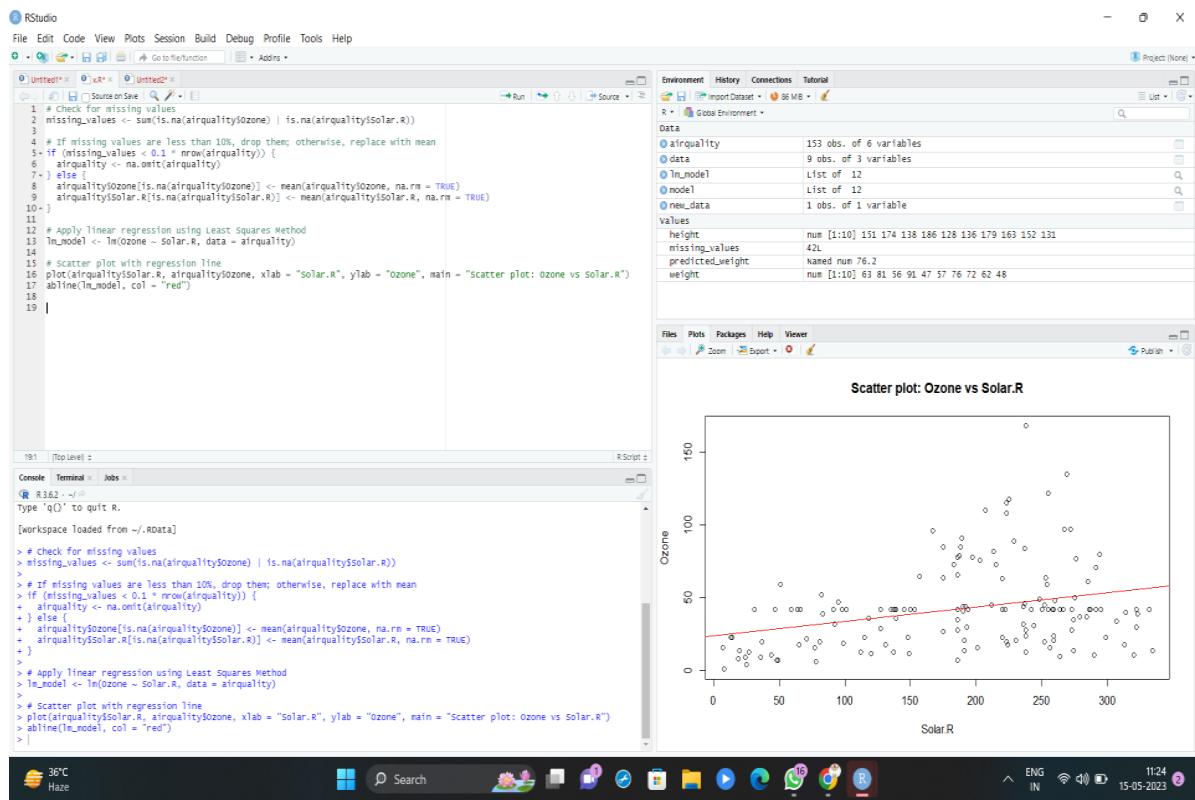
```

main.r
1 summary_stats <- summary(airquality)
2 print(summary_stats)
3 library(reshape2)
4 melted_data <- melt(airquality)
5 print(melted_data)
6 melted_data_id <- melt(airquality, id.vars = c("Month", "Day"))
7 print(melted_data_id)
8 casted_data <- dcast(melted_data_id, Month + Day ~ variable)
9 print(casted_data)
10 averages <- cast(melted_data, Month ~ variable, mean)
11 print(averages)
12

```

The 'Run' button is highlighted in blue. To the right, the 'Output' panel displays the results of the R script execution. It includes a command line prompt 'Rscript /tmp/ITI5SSvPf.r', followed by statistical summaries for 'Ozone', 'Solar.R', 'Wind', and 'Temp' variables, and a summary for the 'Month' and 'Day' factors.

- 4.(i) Find any missing values(na) in features and drop the missing values if its less than 10% else replace that with mean of that feature.
(ii) Apply a linear regression algorithm using Least Squares Method on “Ozone” and “Solar.R”
(iii) Plot Scatter plot between Ozone and Solar and add regression line created by above model



Set-II

1. (i) Write a function to find the factorial of a given number using “for” Loop
- (ii) Create a 3×4 matrix with 12 random numbers between 1-100; have the matrix be filled our row by-row, instead of column-by-column. Name the columns of the matrix *uno*, *dos*, *tres*, *cuatro*, and the rows *x*, *y*, *z*. Scale the matrix by 10 and save the result.
- (iii) Extract the column called “uno” as a vector from the original matrix and save the result

The screenshot shows a web-based R compiler interface from Programiz. The code editor contains the following R script:

```

main.r
1 factorial <- function(n) {
2   result <- 1
3   for (i in 1:n) {
4     result <- result * i
5   }
6   return(result)
7 }
8 set.seed(42) # For reproducibility
9 matrix_data <- matrix(sample(1:100, 12), nrow = 3, ncol = 4, byrow = TRUE,
10   dimnames = list(c("x", "y", "z"), c("uno", "dos", "tres", "cuatro")))
11 scaled_matrix <- matrix_data * 10
12 column_uno <- matrix_data[, "uno"]
13 print(column_uno)
14
15

```

The output window displays the results of the R script:

```

Rscript /tmp/ITI5SSvPf.r
uno dos tres cuatro
x 490 650 250 740
y 180 1000 470 240
z 710 890 370 200
x y z
49 18 71

```

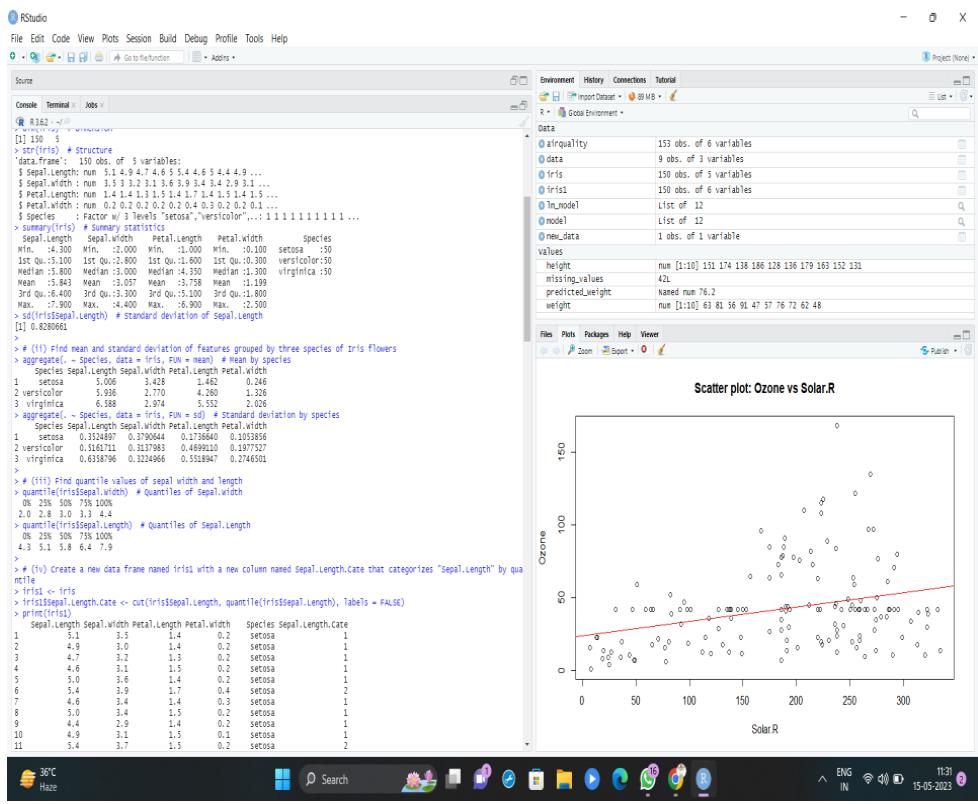
The interface includes a toolbar with icons for file operations, a Run button, and an Output tab. The status bar at the bottom shows system information like temperature (36°C Haze), date (15-05-2023), and time (11:26).

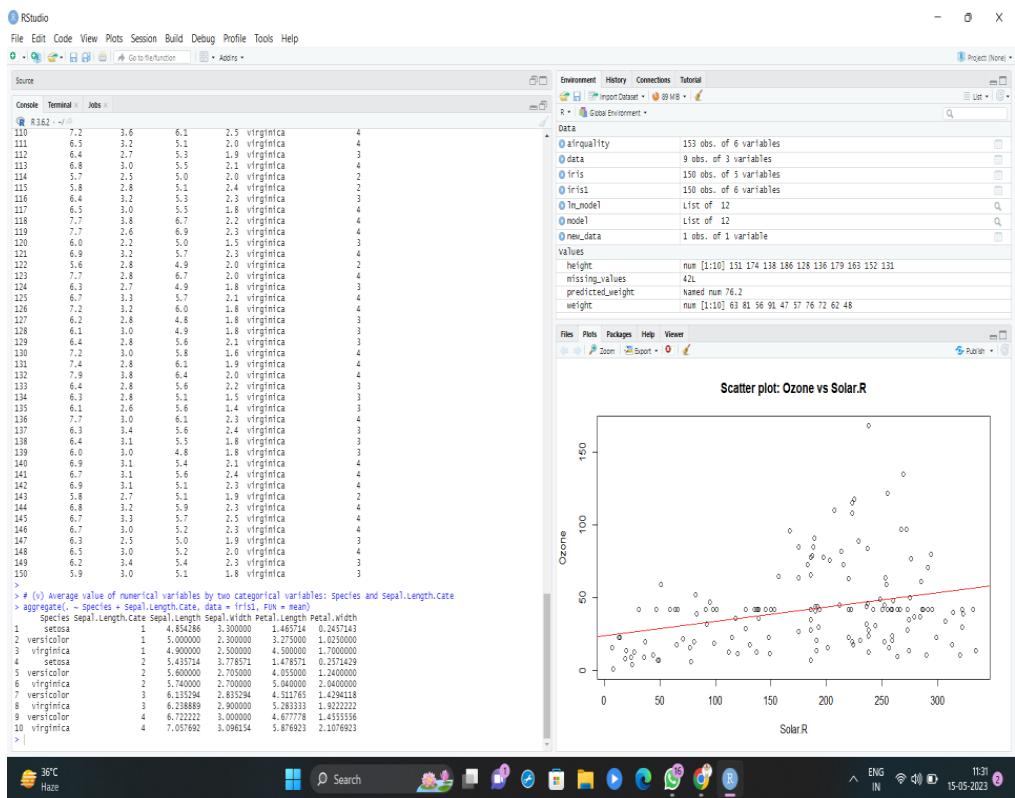
2. In 1936, Edgar Anderson collected data to quantify the geographic variations of iris flowers. The data set consists of 50 samples from each of the three sub-species (*iris setosa*, *iris virginica*, and *iris versicolor*). Four features were measured in centimeters (cm): the lengths and the widths of both sepals and petals

- (i) Find dimension, Structure, Summary statistics, Standard Deviation of all features.
- (ii) Find mean and standard deviation of features groped by three species of Iris flowers (Iris setosa, Iris virginica and Iris versicolor)
- (iii) Find quantile value of sepal width and length
- (iv) create new data frame named iris1 which have a new column name **Sepal.Length.Cate** that categorizes “Sepal.Length” by quantile
- (v) Average value of numerical varialbes by two categorical variables: Species and Sepal.Length.Cate.

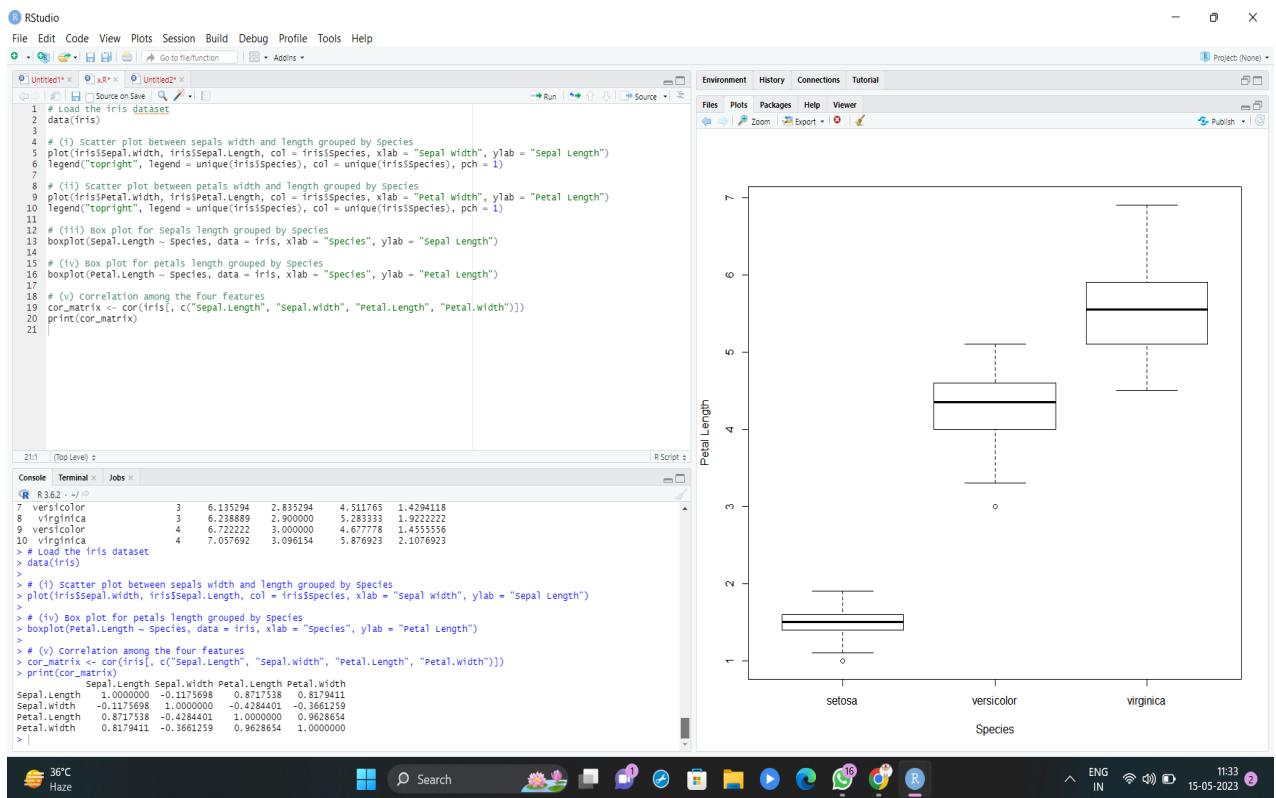
The screenshot shows the RStudio interface with an R script open in the top-left pane. The script contains code for loading the Iris dataset, calculating summary statistics, creating a new categorical variable 'Sepal.Length.Cat', and printing the results. The top-right pane shows the Global Environment, listing objects like 'airquality', 'data', 'iris', and 'iris1'. The bottom-right pane shows the system tray with a weather icon (36°C Haze), search bar, and system status.

```
1 # Load the Iris dataset
2 data(iris)
3
4 # (i) Find dimension, structure, summary statistics, and standard deviation of all features
5 dim(iris) # Dimension
6 str(iris) # Structure
7 summary(iris) # Summary statistics
8 sd(iris$Sepal.Length) # Standard deviation of Sepal.Length
9
10 # (ii) Find mean and standard deviation of features grouped by three species of Iris flowers
11 aggregate(~ Species, data = iris, FUN = mean) # Mean by species
12 aggregate(~ Species, data = iris, FUN = sd) # Standard deviation by species
13
14 # (iii) Find quantile values of sepal width and length
15 quantile(iris$Sepal.Width) # Quantiles of Sepal.Width
16 quantile(iris$Sepal.Length) # Quantiles of Sepal.Length
17
18 # (iv) Create a new data frame named 'iris1' with a new column named 'Sepal.Length.Cat' that categorizes "Sepal.Length" by
19 # species
20 iris1 <- iris
21 iris1$Sepal.Length.Cat <- cut(iris$Sepal.Length, quantile(iris$Sepal.Length), labels = FALSE)
22 print(iris1)
23
24 # (v) Average value of numerical variables by two categorical variables: species and Sepal.Length.Cat
25 aggregate(~ species + Sepal.Length.Cat, data = iris1, FUN = mean)
```





- Plot Scatter plot between sepals width and length grouped by Species
- Plot Scatter plot between petals width and length grouped by Species
- Draw the Box plot for Sepals length grouped by Species
- Draw the Box plot for petals length grouped by Species
- Find the correlation among the four features



- 4.(i) Randomly Sample the iris dataset such as 50% data for training and 50% for test (ii)find summary statistics of above train and test dataset.
 (iii)Create Logistics regression with train data
 (iv)Predict the probability of the model using test data
 (v)Create Confusion matrix for above test model

The screenshot shows a web-based R online compiler interface. On the left, there's a file browser with a file named 'main.r' selected. The code in 'main.r' is as follows:

```
1 data(iris)
2 set.seed(42)
3 train_indices <- sample(1:nrow(iris), nrow(iris) * 0.5)
4 train_data <- iris[train_indices, ]
5 test_data <- iris[-train_indices, ]
6 summary_train <- summary(train_data)
7 summary_test <- summary(test_data)
8 print(summary_train)
9 print(summary_test)
10 logistic_model <- glm(Species ~ ., data = train_data, family = "binomial")
11 predicted_probabilities <- predict(logistic_model, newdata = test_data, type =
  "response")
12 predicted_classes <- ifelse(predicted_probabilities > 0.5, "virginica",
  "versicolor")
13 confusion_matrix <- table(predicted_classes, test_data$Species)
14 print(confusion_matrix)
15
```

The 'Run' button is highlighted in blue. To the right, the 'Output' pane displays the results of the R script execution. It includes descriptive statistics for Sepal.Length, Sepal.Width, Petal.Length, and Petal.Width across three species: setosa, versicolor, and virginica. Below this, a confusion matrix is shown:

| | setosa | versicolor | virginica |
|------------|--------|------------|-----------|
| setosa | 30 | 0 | 0 |
| versicolor | 0 | 21 | 0 |
| virginica | 0 | 0 | 24 |

This screenshot shows the same R online compiler interface as the previous one, but with a different output. The code in 'main.r' is identical.

The 'Run' button is highlighted in blue. The 'Output' pane shows the results of the R script execution. It includes descriptive statistics for Sepal.Length, Sepal.Width, Petal.Length, and Petal.Width across three species: setosa, versicolor, and virginica. Below this, a confusion matrix is shown:

| | setosa | versicolor | virginica |
|------------|--------|------------|-----------|
| setosa | 30 | 0 | 0 |
| versicolor | 0 | 21 | 0 |
| virginica | 0 | 0 | 24 |

At the bottom of the output, there is a 'Warning messages:' section with two entries:

- 1: glm.fit: algorithm did not converge
- 2: glm.fit: fitted probabilities numerically 0 or 1 occurred

Below the warning messages, the confusion matrix is repeated:

| | setosa | versicolor | virginica |
|------------|--------|------------|-----------|
| versicolor | 20 | 0 | 0 |
| virginica | 0 | 29 | 26 |

Set-III

1. Suppose you track your commute times for two weeks (10 days) and you find the following times in minutes 17 16 20 24 22 15 21 15 17 22 Enter this into R as vector data type. (i)create function maxi to find the longest commute time, the function avger to find the average and the function mini to find the minimum.

(ii)Oops, the 24 was a mistake. It should have been 18. How can you fix this? Do so, and then find the new average using above functions.

(iii)How many times was your commute 20 minutes or more?

CODE:

```
# Vector of commute times
commute_times <- c(17, 16, 20, 24, 22, 15, 21, 15, 17, 22)

# Function to find the longest commute time
maxi <- function(commute_times) {
  max(commute_times)
}

# Function to find the average commute time
avger <- function(commute_times) {
  mean(commute_times)
}

# Function to find the minimum commute time
mini <- function(commute_times) {
  min(commute_times)
}

# Testing the functions
longest_commute <- maxi(commute_times)
average_commute <- avger(commute_times)
minimum_commute <- mini(commute_times)

# Printing the results
longest_commute
average_commute
```

```
minimum_commute
```

```
# Correcting the mistake  
commute_times[which(commute_times == 24)] <- 18
```

```
# Recalculating the average using the updated commute times  
new_average_commute <- avgter(commute_times)
```

```
# Printing the new average  
new_average_commute
```

```
# Counting the number of times the commute was 20 minutes or more  
num_greater_than_20 <- sum(commute_times >= 20)
```

```
# Printing the count  
num_greater_than_20
```

The screenshot shows the RGui interface with two windows. The left window is the R Console, displaying the R script and its output. The right window is the R Editor, displaying the results of the script execution.

R Gui (64-bit)

R Console

```
>  
> # Printing the results  
> longest_commute  
[1] 24  
> average_commute  
[1] 18.9  
> minimum_commute  
[1] 15  
>  
> # Correcting the mistake  
> commute_times[which(commute_times == 24)] <- 18  
>  
> # Recalculating the average using the updated commute time  
> new_average_commute <- avgter(commute_times)  
>  
> # Printing the new average  
> new_average_commute  
[1] 18.3  
>  
> # Counting the number of times the commute was 20 minutes  
> num_greater_than_20 <- sum(commute_times >= 20)  
>  
> # Printing the count  
> num_greater_than_20  
[1] 4
```

R C:\Users\BHARATH\OneDrive\Documents\set3-1.R - R Editor

```
# Testing the functions  
longest_commute <- maxi(commute_times)  
average_commute <- avgter(commute_times)  
minimum_commute <- mini(commute_times)  
  
# Printing the results  
longest_commute  
average_commute  
minimum_commute  
  
# Correcting the mistake  
commute_times[which(commute_times == 24)] <- 18  
  
# Recalculating the average using the updated commute times  
new_average_commute <- avgter(commute_times)  
  
# Printing the new average  
new_average_commute  
  
# Counting the number of times the commute was 20 minutes or more  
num_greater_than_20 <- sum(commute_times >= 20)  
  
# Printing the count  
num_greater_than_20
```

2. There is a popular built-in data set in R called "**mtcars**" (Motor Trend Car Road

Tests), which is retrieved from the 1974 Motor Trend US Magazine.

(i)Find the dimension of the data set

(ii)Give the statistical summary of the features.

(iii)Find the largest and smallest value of the variable hp (horsepower).

(iv)Give the mean of mileage per gallon (mpg) with respect to transmission model (feature named as ‘am’)

(v)Give the median of horsepower (hp) with respect to cylinder displacement(cyl)

CODE:

```
# Load the mtcars dataset
data(mtcars)

# Find the dimension of the dataset
dim(mtcars)

# Get the summary statistics of the features
summary(mtcars)
# Find the largest and smallest value of "hp"
max_hp <- max(mtcars$hp)
min_hp <- min(mtcars$hp)

# Print the largest and smallest value of "hp"
max_hp
min_hp
# Calculate the mean of mpg with respect to am
mean_mpg_by_am <- tapply(mtcars$mpg, mtcars$am, mean)

# Print the mean of mpg by am
mean_mpg_by_am
# Calculate the median of hp with respect to cyl
median_hp_by_cyl <- tapply(mtcars$hp, mtcars$cyl, median)

# Print the median of hp by cyl
median_hp_by_cyl
```

```

RGui (64-bit)
File Edit Packages Windows Help
R Console
Max.    :11.0000  Max.   :5.0000  Max.   :8.0000
> # Find the largest and smallest value of "hp"
> max_hp <- max(mtcars$hp)
> min_hp <- min(mtcars$hp)
>
> # Print the largest and smallest value of "hp"
> max_hp
[1] 335
> min_hp
[1] 52
> # Calculate the mean of mpg with respect to am
> mean_mpg_by_am <- tapply(mtcars$mpg, mtcars$am, mean)
>
> # Print the mean of mpg by am
> mean_mpg_by_am
      0        1
17.14737 24.39231
> # Calculate the median of hp with respect to cyl
> median_hp_by_cyl <- tapply(mtcars$hp, mtcars$cyl, median)
>
> # Print the median of hp by cyl
> median_hp_by_cyl
      4       6      8
91.0 110.0 192.5
> |
```

```

C:\Users\BHARATH\OneDrive\Documents\set3-1.R - R Editor
# Load the mtcars dataset
data(mtcars)

# Find the dimension of the dataset
dim(mtcars)

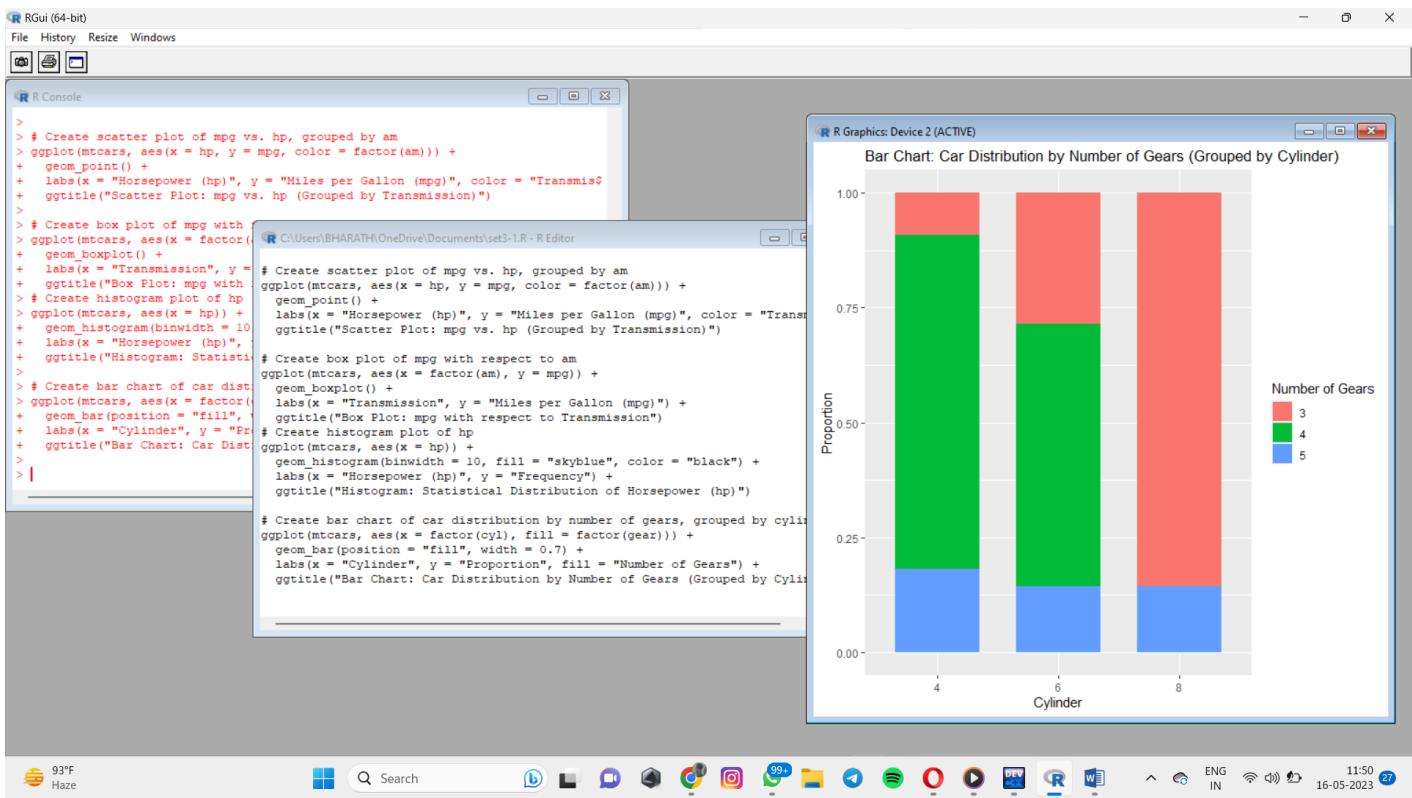
# Get the summary statistics of the features
summary(mtcars)
# Find the largest and smallest value of "hp"
max_hp <- max(mtcars$hp)
min_hp <- min(mtcars$hp)

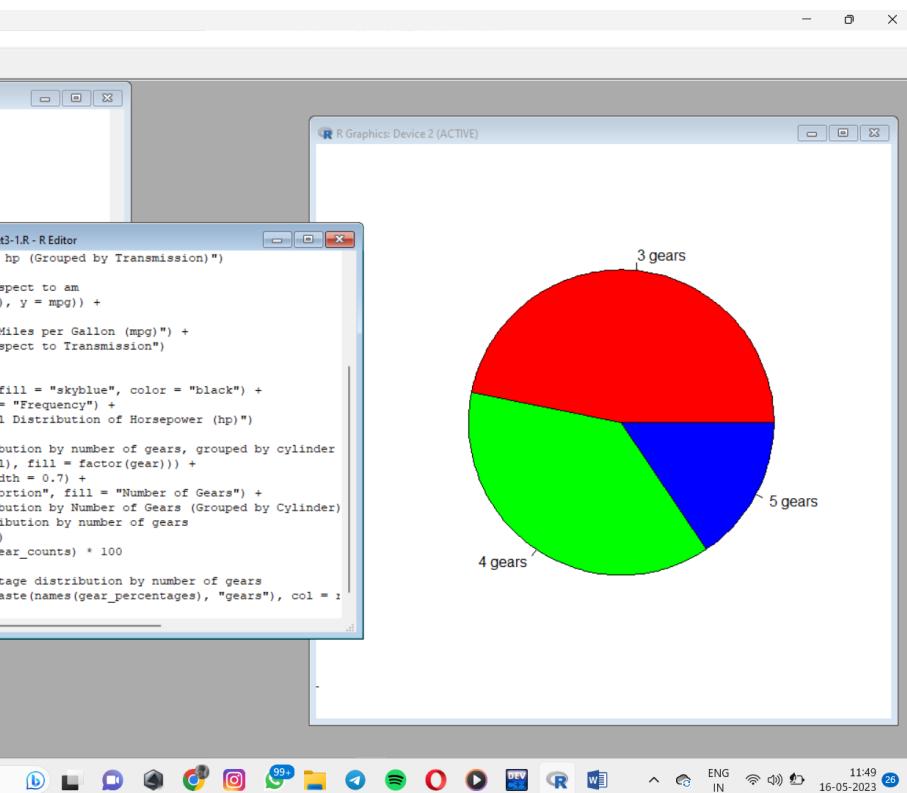
# Print the largest and smallest value of "hp"
max_hp
min_hp
# Calculate the mean of mpg with respect to am
mean_mpg_by_am <- tapply(mtcars$mpg, mtcars$am, mean)

# Print the mean of mpg by am
mean_mpg_by_am
# Calculate the median of hp with respect to cyl
median_hp_by_cyl <- tapply(mtcars$hp, mtcars$cyl, median)

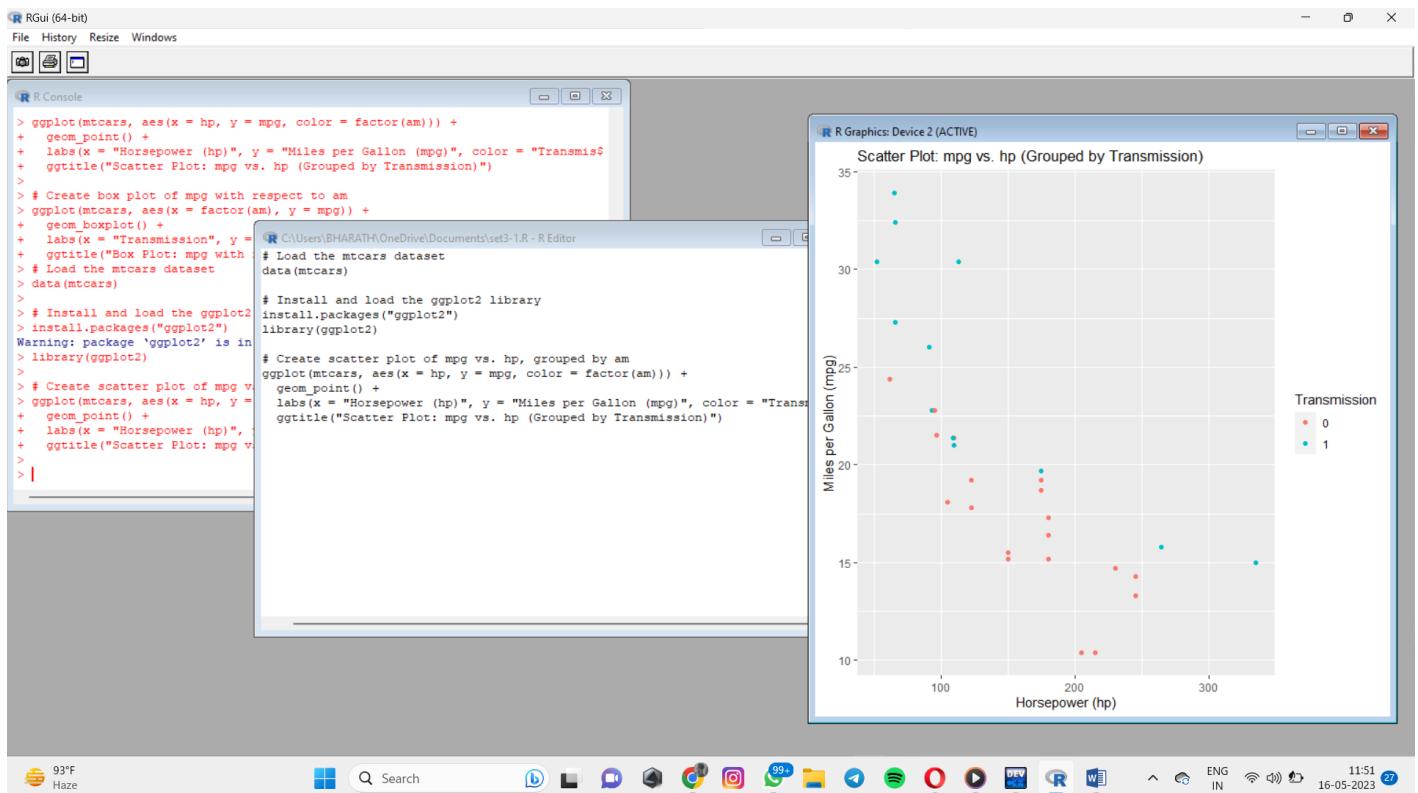
# Print the median of hp by cyl
median_hp_by_cyl
```

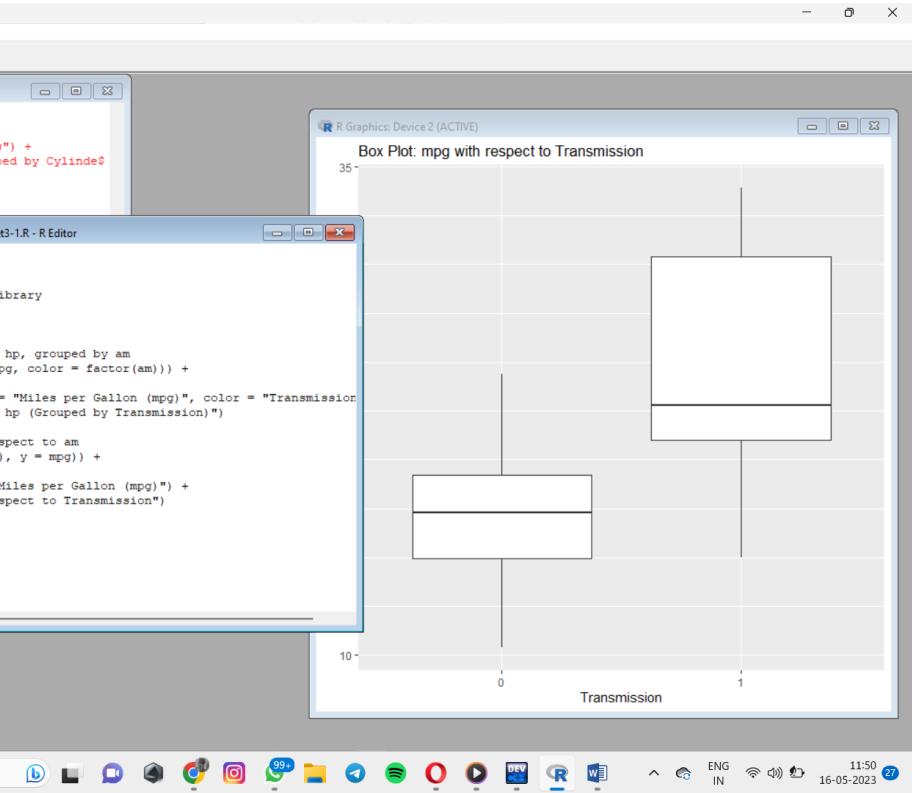
- 3.(i)Create Scatter plot mpg vs hp, grouped by transmission model (feature named as 'am') (ii)Create Box plot for mpg with respect to transmission model (feature named as 'am') (iii)Create histogram plot which shows statistical distribution of hp
(iv)Draw the Bar Chart to show car distribution with respect to number of gears grouped by cylinder.(Grouped or multiple bar chart)
(v)Draw Pie chart which shows the percentage of distribution by number of gears. 4. (i)Generate a multiple regression model using the built-in dataset mtcars. Establish the relationship between "mpg" as a response variable with "disp","hp" and "wt" as predictor variables . (ii)Plot the multiple regression line model with above model parameters.
(iii) Predict the mileage of the car with disp=221, hp=102 and wt=2.91





11:49 16-05-2023 23





Set IV

1. (i) Write a function in R programming to print generate Fibonacci sequence using Recursion in R .
(ii) Find sum of natural numbers up-to 10, without formula using loop statement.
(iii) create a vector 1:10 and Find a square of each number and store that in a separate list.

CODE:

```
fibonacci <- function(n) {  
  if (n <= 1) {  
    return(n)  
  } else {  
    return(fibonacci(n - 1) + fibonacci(n - 2))  
  }  
}
```

```
# Testing the function
fibonacci_sequence <- sapply(0:10, fibonacci)
print(fibonacci_sequence)

# Sum of natural numbers up to 10 using a loop statement
sum_numbers <- 0

for (i in 1:10) {
  sum_numbers <- sum_numbers + i
}

print(sum_numbers)
# Sum of natural numbers up to 10 using a loop statement
sum_numbers <- 0

for (i in 1:10) {
  sum_numbers <- sum_numbers + i
}

print(sum_numbers)
```

The screenshot shows two windows from the RGui interface. On the left is the R Console window, which displays R code and its output. The code includes testing a function for generating a Fibonacci sequence using sapply, and calculating the sum of natural numbers up to 10 using both loop statements and the sapply function. The output shows the sequence [1] 0 1 1 2 3 5 8 13 21 34 55 and the sum [1] 55. On the right is the R Editor window, which contains the same code as the console, likely for comparison or editing.

2. **mtcars**(motor trend car road test) comprises fuel consumption, performance and 10 aspects of automobile design for 32 automobiles. It comes pre-installed with **dplyr** package in R.

- (i)Find the dimension of the data set
- (ii)Give the statistical summary of the features.
- (iii)Print the categorical features in Dataset
- (iv)Find the average weight(wt) grouped by Engine shape(vs)
- (v)Find the largest and smallest value of the variable weight with respect to Engine shape

CODE:

```
# Load the dplyr package
```

```
library(dplyr)
```

```
# Find the dimension of the dataset
```

```
dim(mtcars)
```

```
# Get the summary statistics of the features
summary(mtcars)

# Print the categorical features in the dataset
categorical_features <- mtcars %>%
  select_if(is.factor)

print(categorical_features)
# Print the categorical features in the dataset
categorical_features <- mtcars %>%
  select_if(is.factor)

print(categorical_features)
# Find the average weight (wt) grouped by Engine shape (vs)
average_weight_by_engine_shape <- mtcars %>%
  group_by(vs) %>%
  summarize(average_weight = mean(wt))

print(average_weight_by_engine_shape)
# Find the largest and smallest value of weight with respect to Engine shape
largest_weight_by_engine_shape <- mtcars %>%
  group_by(vs) %>%
  filter(wt == max(wt))

smallest_weight_by_engine_shape <- mtcars %>%
  group_by(vs) %>%
  filter(wt == min(wt))

print(largest_weight_by_engine_shape)
print(smallest_weight_by_engine_shape)
```

RGui (64-bit) - [R Console]

File Edit View Misc Packages Windows Help

```
> categorical_features <- mtcars %>%
+   select_if(is.factor)
>
> print(categorical_features)
data frame with 0 columns and 32 rows
> # Print the categorical features in the dataset
> categorical_features <- mtcars %>%
+   select_if(is.factor)
>
> print(categorical_features)
data frame with 0 columns and 32 rows
> # Find the average weight (wt) grouped by Engine shape (vs)
> average_weight_by_engine_shape <- mtcars %>%
+   group_by(vs) %>%
+   summarize(average_weight = mean(wt))
>
> print(average_weight_by_engine_shape)
# A tibble: 2 × 2
  vs     average_weight
  <dbl>        <dbl>
1 0             3.69
2 1             2.61
> # Find the largest and smallest value of weight with respect to Engine shape
> largest_weight_by_engine_shape <- mtcars %>%
+   group_by(vs) %>%
+   filter(wt == max(wt))
>
> smallest_weight_by_engine_shape <- mtcars %>%
+   group_by(vs) %>%
+   filter(wt == min(wt))
>
> print(largest_weight_by_engine_shape)
# A tibble: 2 × 11
  Groups:   vs [2]
  mpg cyl disp  hp drat    wt  qsec    vs    am gear carb
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 18.1     6   225 105 2.76 3.46 20.2     1     0     3     1
2 10.4     8   460 215 3   5.42 17.8     0     0     3     4
> print(smallest_weight_by_engine_shape)
# A tibble: 2 × 11
  Groups:   vs [2]
  mpg cyl disp  hp drat    wt  qsec    vs    am gear carb
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 26      4   120. 91 4.43 2.14 16.7     0     1     5     2
```

93°F Haze

Search

12:02 16-05-2023

RGUI (64-bit) - [R Console]

File Edit View Misc Packages Windows Help

intersect, setdiff, setequal, union

>

> # Find the dimension of the dataset

> dim(mtcars)

[1] 32 11

> # Get the summary statistics of the features

> summary(mtcars)

| | mpg | cyl | disp | hp |
|---------|--------|---------------|---------------|---------------|
| Min. | :10.40 | Min. :4.000 | Min. :71.1 | Min. :52.0 |
| 1st Qu. | :15.43 | 1st Qu.:4.000 | 1st Qu.:120.8 | 1st Qu.: 96.5 |
| Median | :19.20 | Median :6.000 | Median :196.3 | Median :123.0 |
| Mean | :20.09 | Mean :6.188 | Mean :230.7 | Mean :146.7 |
| 3rd Qu. | :22.80 | 3rd Qu.:8.000 | 3rd Qu.:326.0 | 3rd Qu.:180.0 |
| Max. | :33.90 | Max. :8.000 | Max. :472.0 | Max. :335.0 |

| | drat | wt | qsec | vs |
|---------|--------|---------------|---------------|-----------------|
| Min. | :2.760 | Min. :1.513 | Min. :14.50 | Min. : 0.0000 |
| 1st Qu. | :3.080 | 1st Qu.:2.581 | 1st Qu.:16.89 | 1st Qu.: 0.0000 |
| Median | :3.695 | Median :3.325 | Median :17.71 | Median : 0.0000 |
| Mean | :3.597 | Mean :3.217 | Mean :17.85 | Mean : 0.4375 |
| 3rd Qu. | :3.920 | 3rd Qu.:3.610 | 3rd Qu.:18.90 | 3rd Qu.: 1.0000 |
| Max. | :4.930 | Max. :5.424 | Max. :22.90 | Max. : 1.0000 |

| | am | gear | carb |
|---------|---------|---------------|---------------|
| Min. | :0.0000 | Min. :3.000 | Min. : 1.000 |
| 1st Qu. | :0.0000 | 1st Qu.:3.000 | 1st Qu.:2.000 |
| Median | :0.0000 | Median :4.000 | Median :2.000 |
| Mean | :0.4062 | Mean :3.688 | Mean :2.812 |
| 3rd Qu. | :1.0000 | 3rd Qu.:4.000 | 3rd Qu.:4.000 |
| Max. | :1.0000 | Max. :5.000 | Max. :8.000 |

> # Print the categorical features in the dataset

> categorical_features <- mtcars %>%

+ select_if(is.factor)

>

> print(categorical_features)

data frame with 0 columns and 32 rows

> # Print the categorical features in the dataset

> categorical_features <- mtcars %>%

+ select_if(is.factor)

>

> print(categorical_features)

data frame with 0 columns and 32 rows

> # Find the average weight (wt) grouped by Engine shape (vs)

> average_weight_by_engine_shape <- mtcars %>%

+ group_by(vs) %>%

3. Use ggplot package to plot below EDA questions label the plot accordingly (i) Create weight(wt) vs displacement(disp) scatter plot factor by Engine Shape(vs) (ii) Create horsepower (hp) vs mileage (mpg) scatter plot factor by Engine Shape(vs) (iv) In above(ii) plot , Separate columns according to cylinders(cyl) size
(v) Create histogram plot for horsepower (hp) with bin-width size of 5

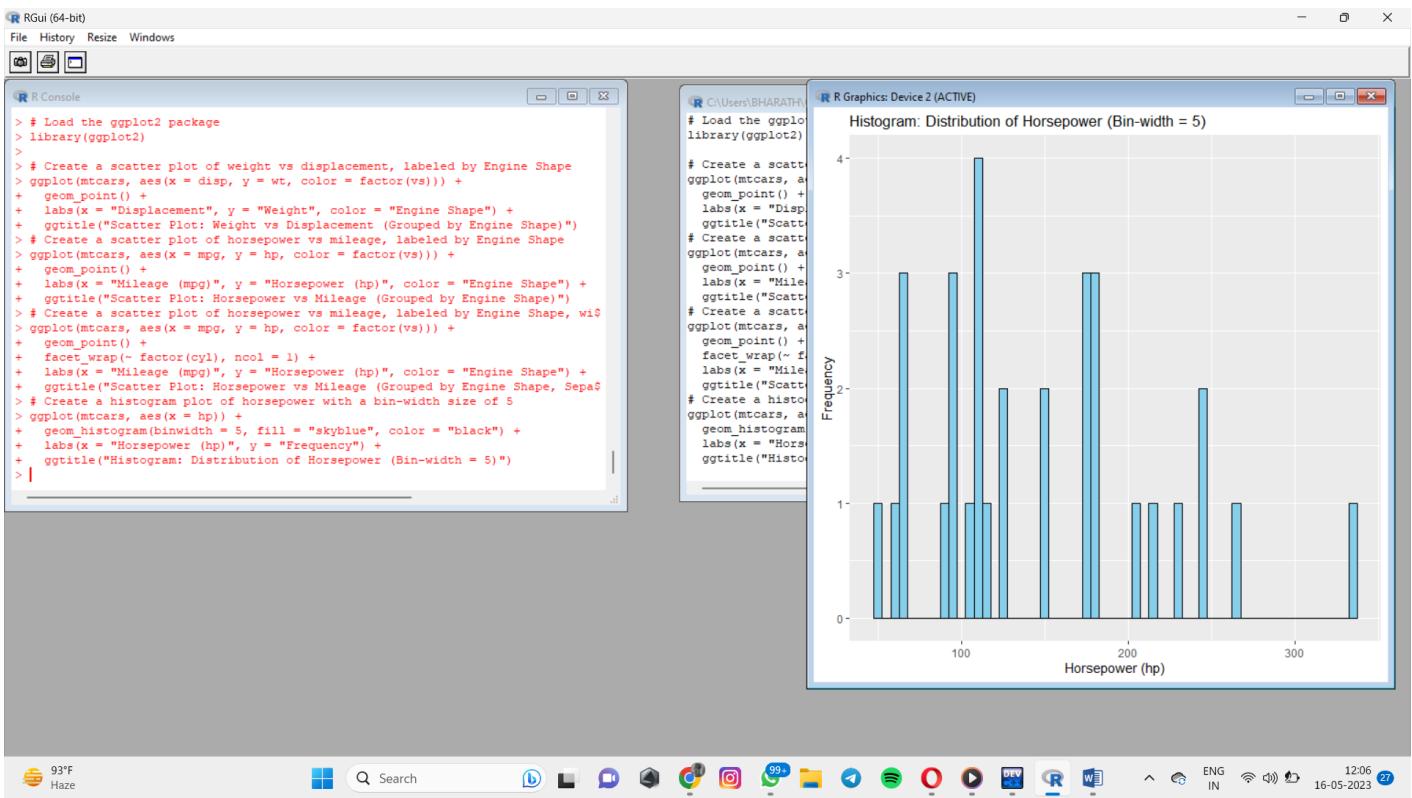
```
# Load the ggplot2 package
library(ggplot2)

# Create a scatter plot of weight vs displacement, labeled by Engine Shape
ggplot(mtcars, aes(x = disp, y = wt, color = factor(vs))) +
  geom_point() +
  labs(x = "Displacement", y = "Weight", color = "Engine Shape") +
  ggtitle("Scatter Plot: Weight vs Displacement (Grouped by Engine Shape)")
```

```
# Create a scatter plot of horsepower vs mileage, labeled by Engine Shape
ggplot(mtcars, aes(x = mpg, y = hp, color = factor(vs))) +
  geom_point() +
  labs(x = "Mileage (mpg)", y = "Horsepower (hp)", color = "Engine Shape") +
  ggtitle("Scatter Plot: Horsepower vs Mileage (Grouped by Engine Shape)")

# Create a scatter plot of horsepower vs mileage, labeled by Engine Shape, with separate columns for
# cylinder size
ggplot(mtcars, aes(x = mpg, y = hp, color = factor(vs))) +
  geom_point() +
  facet_wrap(~ factor(cyl), ncol = 1) +
  labs(x = "Mileage (mpg)", y = "Horsepower (hp)", color = "Engine Shape") +
  ggtitle("Scatter Plot: Horsepower vs Mileage (Grouped by Engine Shape, Separate Columns for Cylinder
Size)")

# Create a histogram plot of horsepower with a bin-width size of 5
ggplot(mtcars, aes(x = hp)) +
  geom_histogram(binwidth = 5, fill = "skyblue", color = "black") +
  labs(x = "Horsepower (hp)", y = "Frequency") +
  ggtitle("Histogram: Distribution of Horsepower (Bin-width = 5)")
```



4. Performing Logistic regression on dataset to predict the cars Engine shape(vs) . (i)Do the EDA analysis and find the features which is impact the Engine shape and use this for model.

- (ii) Split the data set randomly with 80:20 ration to create train and test dataset and create logistic model
- (iii)Create the Confusion matrix among prediction and test data.

RGui (64-bit)

File Edit Packages Windows Help

R Console

```
# Now:
data %>% select(all_of(selected_features))

See <https://tidyselect.r-lib.org/reference/faq-external-vector.html>.
This warning is displayed once every 8 hours.
Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
generated.
>
> # Confirm the selected features
> summary(mtcars_selected)
   hp      wt      am      gear
Min. : 52.0  Min. :1.513  Min. :0.0000  Min. :3.0
1st Qu.: 96.5  1st Qu.:2.581  1st Qu.:0.0000  1st Qu.:3.0
Median :123.0  Median :3.325  Median :0.0000  Median :4.0
Mean   :146.7  Mean   :3.217  Mean   :0.4062  Mean   :3.6
3rd Qu.:180.0  3rd Qu.:3.610  3rd Qu.:1.0000  3rd Qu.:4.0
Max.  :335.0  Max.  :5.424  Max.  :1.0000  Max.  :5.0
   vs
Min. :0.0000
1st Qu.:0.0000
Median :0.0000
Mean   :0.4375
3rd Qu.:1.0000
Max.  :1.0000
> |
```

R Untitled - R Editor

```
geom_boxplot() +
  facet_wrap(~Feature, scales = "free_y") +
  labs(x = "Engine Shape (vs)", y = "Value") +
  theme_bw()

# Bar plots of categorical features with respect to Engine Shape
mtcars %>%
  select_if(is.factor) %>%
  gather(key = "Feature", value = "Value") %>%
  count(Feature, Value) %>%
  ggplot(aes(x = Feature, y = n, fill = Value)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(x = "Feature", y = "Count") +
  theme_bw()

# Based on the EDA analysis, select the features that seem to impact Engine Shape
selected_features <- c("hp", "wt", "am", "gear")

# Create a new dataframe with the selected features
mtcars_selected <- mtcars %>%
  select(selected_features, vs)

# Confirm the selected features
summary(mtcars_selected)
```

R Graphics...

93°F Haze

Search

12:17 16-05-2023

RGui (64-bit)

File Edit Packages Windows Help

R Console

```
Call:
glm(formula = vs ~ ., family = binomial, data = trainData)

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-1.97049 -0.01437  0.00951  0.11610  1.49829

Coefficients:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) 23.90863  24.89812   0.960   0.337    
hp          -0.11359   0.07411  -1.533   0.125    
wt         -1.76182   3.02408  -0.583   0.560    
am        -4.72963   3.77323  -1.253   0.210    
gear       -0.65705   2.56202  -0.256   0.798    

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 36.044 on 25 degrees of freedom
Residual deviance: 10.516 on 21 degrees of freedom
AIC: 20.516

Number of Fisher Scoring iterations: 9
```

> |

R Untitled - R Editor

```
selected_features <- c("hp", "wt", "am", "gear")

# Create a new dataframe with the selected features
mtcars_selected <- mtcars %>%
  select(selected_features, vs)

# Confirm the selected features
summary(mtcars_selected)
# Load the necessary libraries
library(caret)

# Set the seed for reproducibility
set.seed(123)

# Split the dataset into training and testing sets
trainIndex <- createDataPartition(mtcars_selected$vs, p = 0.8, list = FALSE)
trainData <- mtcars_selected[trainIndex, ]
testData <- mtcars_selected[-trainIndex, ]

# Create a logistic regression model
logistic_model <- glm(vs ~ ., data = trainData, family = binomial)

# Print the model summary
summary(logistic_model)
```

R Graphics...

93°F Haze

Search

12:18 16-05-2023

The screenshot shows the RGui interface with two windows open:

- R Console**: Displays R code and its output. The output includes:


```
(Dispersion parameter for binomial family taken to be 1)

Null deviance: 36.044 on 25 degrees of freedom
Residual deviance: 10.516 on 21 degrees of freedom
AIC: 20.516

Number of Fisher Scoring iterations: 9

> # Make predictions using the logistic regression model
> predictions <- predict(logistic_model, newdata = testData,
>
> # Convert predicted probabilities to predicted classes
> predicted_classes <- ifelse(predictions >= 0.5, 1, 0)
>
> # Create the confusion matrix
> confusion_matrix <- table(predicted_classes, testData$vs)
>
> # Print the confusion matrix
> print(confusion_matrix)

predicted_classes 0 1
                 0 5 0
                 1 0 1
> |
```
- R Editor**: Displays R code for a logistic regression model. The code includes:


```
# Set the seed for reproducibility
set.seed(123)

# Split the dataset into training and testing sets
trainIndex <- createDataPartition(mtcars_selected$vs, p = 0.8, list = FALSE)
trainData <- mtcars_selected[trainIndex, ]
testData <- mtcars_selected[-trainIndex, ]

# Create a logistic regression model
logistic_model <- glm(vs ~ ., data = trainData, family = binomial)

# Print the model summary
summary(logistic_model)
# Make predictions using the logistic regression model
predictions <- predict(logistic_model, newdata = testData, type = "response")

# Convert predicted probabilities to predicted classes
predicted_classes <- ifelse(predictions >= 0.5, 1, 0)

# Create the confusion matrix
confusion_matrix <- table(predicted_classes, testData$vs)

# Print the confusion matrix
print(confusion_matrix)
```

Set-V

1.(i) Write a R program to extract the five of the levels of factor created from a random sample from the LETTERS (Part of the base R distribution.)

(ii) Write R function to find the range of given vector. Range=Max-Min

Sample input, C<-(9,8,7,6,5,4,3,2,1), output=8

(iii) Write the R function to find the number of vowels in given string

Sample input c<- “matrix”, output<-2

CODE:

```
# Set the seed for reproducibility
```

```
set.seed(123)
```

```
# Create a random sample from the LETTERS vector
random_sample <- sample(LETTERS, size = 20, replace = TRUE)

# Convert the sample into a factor
factor_sample <- as.factor(random_sample)

# Extract the first five levels of the factor
extracted_levels <- levels(factor_sample)[1:5]

# Print the extracted levels
print(extracted_levels)

# Function to find the range of a given vector
find_range <- function(vector) {
  range_value <- max(vector) - min(vector)
  return(range_value)
}

# Sample input vector
C <- c(9, 8, 7, 6, 5, 4, 3, 2, 1)

# Find the range of the input vector using the function
output_range <- find_range(C)
```

```
# Print the output range
print(output_range)

# Function to find the number of vowels in a given string
count_vowels <- function(string) {
  vowels <- c("a", "e", "i", "o", "u")
  count <- sum(strsplit(tolower(string), "")[[1]] %in% vowels)
  return(count)
}

# Sample input string
input_string <- "Hello, how are you?"

# Find the number of vowels in the input string using the function
output_count <- count_vowels(input_string)

# Print the output count
print(output_count)
```

The screenshot shows the RGui interface with the title bar "RGui (64-bit) - [R Console]". The menu bar includes File, Edit, View, Misc, Packages, Windows, and Help. Below the menu is a toolbar with various icons. The main area contains R code and its output:

```
> # Create a random sample from the LETTERS vector
> random_sample <- sample(LETTERS, size = 20, replace = TRUE)
>
> # Convert the sample into a factor
> factor_sample <- as.factor(random_sample)
>
> # Extract the first five levels of the factor
> extracted_levels <- levels(factor_sample)[1:5]
>
> # Print the extracted levels
> print(extracted_levels)
[1] "C" "E" "I" "J" "K"
> # Function to find the range of a given vector
> find_range <- function(vector) {
+   range_value <- max(vector) - min(vector)
+   return(range_value)
+ }
>
> # Sample input vector
> C <- c(9, 8, 7, 6, 5, 4, 3, 2, 1)
>
> # Find the range of the input vector using the function
> output_range <- find_range(C)
>
> # Print the output range
> print(output_range)
[1] 8
> # Function to find the number of vowels in a given string
> count_vowels <- function(string) {
+   vowels <- c("a", "e", "i", "o", "u")
+   count <- sum(strsplit(tolower(string), "")[[1]] %in% vowels)
+   return(count)
+ }
>
> # Sample input string
> input_string <- "Hello, how are you?"
>
> # Find the number of vowels in the input string using the function
> output_count <- count_vowels(input_string)
>
> # Print the output count
> print(output_count)
[1] 7
>
```

The taskbar at the bottom shows various application icons, and the system tray indicates it's 12:22 on 16-05-2023.

2. Load inbuild dataset “ChickWeight” in R

- (i) Explore the summary of Data set, like number of Features and its type. Finds the number of records for each features
- (ii) Extract last 6 records of dataset
- (iii) order the data frame, in ascending order by feature name “weight” grouped by feature “diet”
- (iv) Perform melting function based on “Chick”, “Time”, “Diet” features as ID variables
- (v) Perform cast function to display the mean value of weight grouped by Diet

The screenshot shows the RGui (64-bit) interface with three main windows:

- R Console**: Displays R code and its output. The code loads the ChickWeight dataset, summarizes it, and prints feature types and record counts. The output shows 4 features, 3 feature types (numeric, ordered, factor), and 578 records per feature.
- R Editor**: Shows the same R code as the console, intended for editing or copying and pasting.
- R Graphics...**: A placeholder window for displaying plots.

The taskbar at the bottom includes icons for various applications like File Explorer, Edge, and Spotify, along with system status indicators for battery, signal, and date/time (16-05-2023, 12:25).

```
>
> # Print the results
> print(paste("Number of features:", num_features))
[1] "Number of features: 4"
> print("Feature types:")
[1] "Feature types:"
> print(feature_types)
$weight
[1] "numeric"

$Time
[1] "numeric"

$Chick
[1] "ordered" "factor"

$Diet
[1] "factor"

> print("Number of records per feature:")
[1] "Number of records per feature:"
> print(records_per_feature)
      weight   Time   Chick   Diet
  578     578     578     578
> |
```

```
# Load the ChickWeight dataset
data(ChickWeight)

# Check the summary of the dataset
summary(ChickWeight)

# Get the number of features and their types
num_features <- ncol(ChickWeight)
feature_types <- sapply(ChickWeight, class)

# Get the number of records for each feature
records_per_feature <- sapply(ChickWeight, length)

# Print the results
print(paste("Number of features:", num_features))
print("Feature types:")
print(feature_types)
print("Number of records per feature:")
print(records_per_feature)
```

The screenshot shows the RGui (64-bit) interface with two windows open:

- R Console**: Displays R code and its output. The code reads the ChickWeight dataset, prints feature types, and extracts the last six records.
- R Editor**: Displays the same R code, likely a copy-paste from the console or a script file.

R Graphics window is also visible at the bottom left.

```
R Gui (64-bit)
File Edit Packages Windows Help
R Console
$Chick
[1] "ordered" "factor"

$Diet
[1] "factor"

> print("Number of records per feature:")
[1] "Number of records per feature:"
> print(records_per_feature)
      weight   Time   Chick   Diet
  578     578     578     578
> # Extract the last 6 records of the dataset
> last_six_records <- tail(ChickWeight, 6)
>
> # Print the extracted records
> print(last_six_records)
Grouped Data: weight ~ Time | Chick
      weight   Time   Chick   Diet
  573     155     12     50     4
  574     175     14     50     4
  575     205     16     50     4
  576     234     18     50     4
  577     264     20     50     4
  578     264     21     50     4
> |
```

```
R Untitled - R Editor
# Load the ChickWeight dataset
data(ChickWeight)

# Check the summary of the dataset
summary(ChickWeight)

# Get the number of features and their types
num_features <- ncol(ChickWeight)
feature_types <- sapply(ChickWeight, class)

# Get the number of records for each feature
records_per_feature <- sapply(ChickWeight, length)

# Print the results
print(paste("Number of features:", num_features))
print("Feature types:")
print(feature_types)
print("Number of records per feature:")
print(records_per_feature)
# Extract the last 6 records of the dataset
last_six_records <- tail(ChickWeight, 6)

# Print the extracted records
print(last_six_records)
```

RGui (64-bit) - [R Console]

File Edit View Misc Packages Windows Help

| 396 | 238 | 14 | 35 | 3 |
|-----|-----|----|----|---|
| 530 | 238 | 21 | 46 | 4 |
| 228 | 240 | 14 | 21 | 2 |
| 165 | 248 | 18 | 14 | 1 |
| 82 | 250 | 18 | 7 | 1 |
| 447 | 250 | 20 | 39 | 3 |
| 292 | 251 | 21 | 26 | 2 |
| 352 | 256 | 21 | 31 | 3 |
| 166 | 259 | 20 | 14 | 1 |
| 279 | 259 | 20 | 25 | 2 |
| 552 | 261 | 18 | 48 | 4 |
| 458 | 262 | 18 | 40 | 3 |
| 362 | 263 | 18 | 32 | 3 |
| 577 | 264 | 20 | 50 | 4 |
| 578 | 264 | 21 | 50 | 4 |
| 280 | 265 | 21 | 25 | 2 |
| 167 | 266 | 21 | 14 | 1 |
| 483 | 269 | 20 | 42 | 4 |
| 448 | 272 | 21 | 39 | 3 |
| 229 | 275 | 16 | 21 | 2 |
| 327 | 279 | 20 | 29 | 2 |
| 435 | 280 | 20 | 38 | 3 |
| 484 | 281 | 21 | 42 | 4 |
| 397 | 287 | 16 | 35 | 3 |
| 83 | 288 | 20 | 7 | 1 |
| 436 | 290 | 21 | 38 | 3 |
| 363 | 291 | 20 | 32 | 3 |
| 386 | 294 | 18 | 34 | 3 |
| 459 | 295 | 20 | 40 | 3 |
| 553 | 303 | 20 | 48 | 4 |
| 84 | 305 | 21 | 7 | 1 |
| 364 | 305 | 21 | 32 | 3 |
| 230 | 307 | 18 | 21 | 2 |
| 328 | 309 | 21 | 29 | 2 |
| 231 | 318 | 20 | 21 | 2 |
| 460 | 321 | 21 | 40 | 3 |
| 554 | 322 | 21 | 48 | 4 |
| 387 | 327 | 20 | 34 | 3 |
| 232 | 331 | 21 | 21 | 2 |
| 398 | 332 | 18 | 35 | 3 |
| 388 | 341 | 21 | 34 | 3 |
| 359 | 361 | 20 | 35 | 3 |
| 400 | 373 | 21 | 35 | 3 |

> |

93°F Haze

Search

12:26 16-05-2023 28

RGui (64-bit)

File Edit Packages Windows Help

R Console

```
556 49 2 4 weight 53
557 49 4 4 weight 64
558 49 6 4 weight 85
559 49 8 4 weight 108
560 49 10 4 weight 128
561 49 12 4 weight 152
562 49 14 4 weight 166
563 49 16 4 weight 184
564 49 18 4 weight 203
565 49 20 4 weight 233
566 49 21 4 weight 237
567 50 0 4 weight 41
568 50 2 4 weight 54
569 50 4 4 weight 67
570 50 6 4 weight 84
571 50 8 4 weight 105
572 50 10 4 weight 122
573 50 12 4 weight 155
574 50 14 4 weight 175
575 50 16 4 weight 205
576 50 18 4 weight 234
577 50 20 4 weight 264
578 50 21 4 weight 264
```

> |

R Graphics...

Untitled - R Editor

```
# Print the results
print(paste("Number of features:", num_features))
print("Feature types:")
print(feature_types)
print("Number of records per feature:")
print(records_per_feature)
# Extract the last 6 records of the dataset
last_six_records <- tail(ChickWeight, 6)

# Print the extracted records
print(last_six_records)
# Order the data frame in ascending order by "weight" grouped by "diet"
ordered_df <- ChickWeight[order(ChickWeight$weight), ]

# Print the ordered data frame
print(ordered_df)
# Perform the melting function
library(reshape2)

melted_df <- melt(ChickWeight, id.vars = c("Chick", "Time", "Diet"))

# Print the melted data frame
print(melted_df)
```

93°F Haze

Search

12:27 16-05-2023 28

The screenshot shows the RGui interface with three main windows:

- R Console**: Displays the following R code and its output. The output shows data for chick 566 to 578, grouped by Diet (1, 2, 3, 4) and Time (4, 8, 10, 12, 14, 16, 18, 20, 21).

```

566 49 21 4 weight 237
567 50 0 4 weight 41
568 50 2 4 weight 54
569 50 4 4 weight 67
570 50 6 4 weight 84
571 50 8 4 weight 105
572 50 10 4 weight 122
573 50 12 4 weight 155
574 50 14 4 weight 175
575 50 16 4 weight 205
576 50 18 4 weight 234
577 50 20 4 weight 264
578 50 21 4 weight 264
>
> # Perform the cast function to get the mean value of "weight"
> cast_df <- dcast(melted_df, Diet ~ variable, mean)
>
> # Print the casted data frame
> print(cast_df)
  Diet weight
1   1 102.6455
2   2 122.6167
3   3 142.9500
4   4 135.2627
> |
```

- R Editor**: Shows the R code used to manipulate the ChickWeight dataset. It includes extracting the last 6 records, ordering the data frame by weight for each diet, melting the data, and calculating the mean weight for each diet.

```

# Extract the last 6 records of the dataset
last_six_records <- tail(ChickWeight, 6)

# Print the extracted records
print(last_six_records)
# Order the data frame in ascending order by "weight" grouped by "diet"
ordered_df <- ChickWeight[order(ChickWeight$weight), ]

# Print the ordered data frame
print(ordered_df)
# Perform the melting function
library(reshape2)

melted_df <- melt(ChickWeight, id.vars = c("Chick", "Time", "Diet"))

# Print the melted data frame
print(melted_df)

# Perform the cast function to get the mean value of "weight" grouped by "Diet"
cast_df <- dcast(melted_df, Diet ~ variable, mean)

# Print the casted data frame
print(cast_df)
```

- R Graphics**: Shows the Windows taskbar at the bottom with various application icons.

3.(i)Get the Statistical Summary of “ChickWeight” dataset

(ii)Create Box plot for “weight” grouped by “Diet”

(iii)Create a Histogram for “Weight” features belong to Diet- 1 category

(iv) Create a Histogram for “Weight” features belong to Diet- 4 category

(v) Create Scatter plot for weight vs Time grouped by Diet

4.(i) Create multi regression model to find a weight of the chicken , by “Time” and “Diet” as as predictor variables

(ii) Predict weight for Time=10 and Diet=1

(iii)Find the error(MAE) in model for same